

Scoring overhead vs attention kernels

